

1 VM:

```
URL: spark://master:7077
Alive Workers: 1
Cores in use: 8 Total, 0 Used
Memory in use: 4.0 GiB Total, 0.0 B Used
Resources in use:
Applications: 0 Running, 1 Completed
Drivers: 0 Running, 0 Completed
Status: ALIVE
```

▼ Workers (1)

Worker Id	Address	State	Cores	Memory	Resources
worker-20251109020828-192.168.13.200-34165	192.168.13.200:34165	ALIVE	8 (0 Used)	4.0 GiB (0.0 B Used)	

▼ Running Applications (0)

Application ID	Name	Cores	Memory per Executor	Resources Per Executor	Submitted Time	User	State	Duration
app-20251109122752-0000	sat516_stroke_pipeline	8	1024.0 MiB		2025/11/09 12:27:52	sat3812	FINISHED	54 min

▼ Completed Applications (1)

Application ID	Name	Cores	Memory per Executor	Resources Per Executor	Submitted Time	User	State	Duration
app-20251109122752-0000	sat516_stroke_pipeline	8	1024.0 MiB		2025/11/09 12:27:52	sat3812	FINISHED	54 min

```
[sat3812@master ~]$ $SPARK_HOME/sbin/start-master.sh
$SPARK_HOME/sbin/start-worker.sh spark://master:7077
org.apache.spark.deploy.master.Master running as process 128448. Stop it first.
org.apache.spark.deploy.worker.Worker running as process 123127. Stop it first.
[sat3812@master ~]$ sudo nano stoke_pipeline.py
[sudo] password for sat3812:
[sat3812@master ~]$ sudo nano stroke_pipeline.py
[sat3812@master ~]$ sudo nano stroke_pipeline.py
[sat3812@master ~]$ spark-submit --master spark://master:7077 stroke_pipeline.py
Loading: hdfs://master:9000/user/sat3812/stroke_project/data/healthcare-dataset-stroke-data.csv
Oversampled minority from 191 to ~4011 vs majority 3932.
[info] MLP input_dim = 20

Training Logistic Regression...
Total time: 242.91 seconds

Logistic Regression Performance Results
accuracy: 0.727
auc: 0.861
precision: 0.154
recall: 0.810
specificity: 0.722

Training Linear SVM...
Total time: 1065.26 seconds

Linear SVM Performance Results
accuracy: 0.719
auc: 0.861
precision: 0.155
recall: 0.845
specificity: 0.712

Training Gradient Boosted Trees...
Total time: 1545.94 seconds

Gradient Boosted Trees Performance Results
accuracy: 0.819
auc: 0.802
precision: 0.162
recall: 0.500
specificity: 0.839

Training Multilayer Perceptron...
Total time: 258.94 seconds

Multilayer Perceptron Performance Results
accuracy: 0.781
auc: 0.697
precision: 0.097
recall: 0.328
specificity: 0.809
MODEL COMPARISON SUMMARY
-----
Model          Acc   AUC   Prec  Recall  Spec  Time(s)
-----
Logistic Regression  0.727  0.861  0.154  0.810  0.722  242.91
Linear SVM        0.719  0.861  0.155  0.845  0.712  1065.26
Gradient Boosted Trees  0.819  0.802  0.162  0.500  0.839  1545.94
Multilayer Perceptron  0.781  0.697  0.097  0.328  0.809  258.94
[main] done.
[sat3812@master ~]$
```

2 VMS:

URL: spark://master:7077
Alive Workers: 2
Cores in use: 16 Total, 0 Used
Memory in use: 8.0 GiB Total, 0.0 B Used
Resources in use:
Applications: 0 Running, 1 Completed
Drivers: 0 Running, 0 Completed
Status: ALIVE

▼ Workers (2)

Worker Id	Address	State	Cores	Memory	Resources
worker-20251109020828-192.168.13.200-34165	192.168.13.200:34165	ALIVE	8 (0 Used)	4.0 GiB (0.0 B Used)	
worker-20251109134934-192.168.13.201-37849	192.168.13.201:37849	ALIVE	8 (0 Used)	4.0 GiB (0.0 B Used)	

▼ Running Applications (0)

Application ID	Name	Cores	Memory per Executor	Resources Per Executor	Submitted Time	User	State	Duration

▼ Completed Applications (1)

Application ID	Name	Cores	Memory per Executor	Resources Per Executor	Submitted Time	User	State	Duration
app-20251109135056-0000	sat5165_stroke_pipeline	16	1024.0 MiB		2025/11/09 13:50:56	sat3812	FINISHED	37 min

```
[sat3812@master ~]$ spark-submit --master spark://master:7077 stroke_pipeline.py
Loading: hdfs://master:9000/user/sat3812/stroke_project/data/healthcare-dataset-stroke-data.csv
Oversampled minority from 191 to -4011 vs majority 3932.
[info] MLP input_dim = 20

Training Logistic Regression...
Total time: 193.03 seconds

Logistic Regression Performance Results
accuracy: 0.727
auc: 0.861
precision: 0.154
recall: 0.810
specificity: 0.722

Training Linear SVM...
Total time: 728.84 seconds

Linear SVM Performance Results
accuracy: 0.719
auc: 0.861
precision: 0.155
recall: 0.845
specificity: 0.712

Training Gradient Boosted Trees...
Total time: 1039.26 seconds

Gradient Boosted Trees Performance Results
accuracy: 0.819
auc: 0.882
precision: 0.162
recall: 0.500
specificity: 0.839

Training Multilayer Perceptron...
Total time: 179.01 seconds

Multilayer Perceptron Performance Results
accuracy: 0.781
auc: 0.697
precision: 0.097
recall: 0.328
specificity: 0.809
MODEL COMPARISON SUMMARY
-----
Model          Acc   AUC   Prec  Recall   Spec  Time(s)
-----
Logistic Regression    0.727  0.861  0.154  0.810  0.722  193.03
Linear SVM           0.719  0.861  0.155  0.845  0.712  728.84
Gradient Boosted Trees 0.819  0.882  0.162  0.500  0.839  1039.26
Multilayer Perceptron  0.781  0.697  0.097  0.328  0.809  179.01
[main] done.
[sat3812@master ~]$
```

3 VMS:

▼ Workers (3)

Worker Id	Address	State	Cores	Memory	Resources
worker-20251109161018-192.168.13.201-32833	192.168.13.201:32833	ALIVE	8 (0 Used)	4.0 GiB (0.0 B Used)	
worker-20251109163402-192.168.13.200-37339	192.168.13.200:37339	ALIVE	8 (0 Used)	4.0 GiB (0.0 B Used)	
worker-20251109164312-192.168.13.122-37287	192.168.13.122:37287	ALIVE	8 (0 Used)	4.0 GiB (0.0 B Used)	

▼ Running Applications (0)

Application ID	Name	Cores	Memory per Executor	Resources Per Executor	Submitted Time	User	State	Duration
app-20251109164405-0000	sat5165_stroke_pipeline	24	2.0 GiB		2025/11/09 16:44:05	sat3812	FINISHED	36 min

▼ Completed Applications (1)

Application ID	Name	Cores	Memory per Executor	Resources Per Executor	Submitted Time	User	State	Duration
app-20251109164405-0000	sat5165_stroke_pipeline	24	2.0 GiB		2025/11/09 16:44:05	sat3812	FINISHED	36 min

```

Loading: hdfs://master:9000/user/sat3812/stroke_project/data/healthcare-dataset-stroke-d
Oversampled minority from 191 to ~4011 vs majority 3932.
[info] MLP input_dim = 20

Training Logistic Regression...
Total time: 203.43 seconds

Logistic Regression Performance Results
accuracy: 0.727
auc: 0.861
precision: 0.154
recall: 0.810
specificity: 0.722

Training Linear SVM...
Total time: 697.93 seconds

Linear SVM Performance Results
accuracy: 0.719
auc: 0.861
precision: 0.155
recall: 0.845
specificity: 0.712

Training Gradient Boosted Trees...
Total time: 916.42 seconds

Gradient Boosted Trees Performance Results
accuracy: 0.819
auc: 0.803
precision: 0.162
recall: 0.500
specificity: 0.839

Training Multilayer Perceptron...
Total time: 181.76 seconds

Multilayer Perceptron Performance Results
accuracy: 0.791
auc: 0.696
precision: 0.102
recall: 0.328
specificity: 0.820
MODEL COMPARISON SUMMARY
-----
Model          Acc   AUC   Prec  Recall  Spec  Time(s)
-----
Logistic Regression  0.727  0.861  0.154  0.810  0.722  203.43
Linear SVM        0.719  0.861  0.155  0.845  0.712  697.93
Gradient Boosted Trees  0.819  0.803  0.162  0.500  0.839  916.42
Multilayer Perceptron  0.791  0.696  0.102  0.328  0.820  181.76
Stroke Pipeline Complete.

```

4 VMS:

▼ Workers (4)

Worker Id	Address	State	Cores	Memory	Resources
worker-20251109192507-192.168.13.201-33947	192.168.13.201:33947	ALIVE	8 (0 Used)	4.0 GiB (0.0 B Used)	
worker-20251109192515-192.168.13.122-43307	192.168.13.122:43307	ALIVE	8 (0 Used)	4.0 GiB (0.0 B Used)	
worker-20251109192515-192.168.13.123-42317	192.168.13.123:42317	ALIVE	8 (0 Used)	4.0 GiB (0.0 B Used)	
worker-20251109192516-192.168.13.200-46823	192.168.13.200:46823	ALIVE	8 (0 Used)	4.0 GiB (0.0 B Used)	

▼ Running Applications (0)

Application ID	Name	Cores	Memory per Executor	Resources Per Executor	Submitted Time	User	State	Duration

▼ Completed Applications (1)

Application ID	Name	Cores	Memory per Executor	Resources Per Executor	Submitted Time	User	State	Duration
app-20251109193815-0000	sat5165_stroke_pipeline	32	3.0 GiB		2025/11/09 19:38:15	sat3812	FINISHED	27 min

```

Loading: hdfs://master:9000/user/sat3812/stroke_project/data/healthcare-dataset-stroke-data.csv
Oversampled minority from 191 to ~4011 vs majority 3932.
[info] MLP input_dim = 20

Training Logistic Regression...
Total time: 164.51 seconds

Logistic Regression Performance Results
accuracy: 0.727
auc: 0.861
precision: 0.154
recall: 0.810
specificity: 0.722

Training Linear SVM...
Total time: 548.78 seconds

Linear SVM Performance Results
accuracy: 0.719
auc: 0.861
precision: 0.155
recall: 0.845
specificity: 0.712

Training Gradient Boosted Trees...
Total time: 690.43 seconds

Gradient Boosted Trees Performance Results
accuracy: 0.819
auc: 0.802
precision: 0.162
recall: 0.500
specificity: 0.839

Training Multilayer Perceptron...
Total time: 132.34 seconds

Multilayer Perceptron Performance Results
accuracy: 0.779
auc: 0.698
precision: 0.096
recall: 0.328
specificity: 0.807

MODEL COMPARISON SUMMARY
-----
Model          Acc   AUC   Prec  Recall  Spec  Time(s)
-----
Logistic Regression  0.727  0.861  0.154  0.810  0.722  164.51
Linear SVM        0.719  0.861  0.155  0.845  0.712  548.78
Gradient Boosted Trees  0.819  0.802  0.162  0.500  0.839  690.43
Multilayer Perceptron  0.779  0.698  0.096  0.328  0.807  132.34
Stroke Pipeline Complete.
[sat3812@master ~]$
```